

# Machine Learning with Cloud and Distributed Systems

Investigating the Application of  
Machine Learning Models to Predict  
Resource Usage Trends within  
Distributed Systems

Raphael Leong



# Background Overview



# Cloud and Distributed Systems



Example of Google Cluster within their data center



# Google Cluster Trace

- Dataset that is collected over a period of 29 days
- Outlines the details of around 12,500 machines
- Timeseries data describes cluster usage and job scheduling
- Dataset used in prior works to research cloud and distributed systems



# Motivations

- To further understand the complexities behind cluster scheduling
- To improve workload distribution and optimisation between nodes
- Investigate and demonstrate the potential of machine learning models within this space



# Objectives

- Analyse the Google cluster trace dataset
- Investigate the relationship between cluster machines based on defined characteristics set by prior works
- Explore the possibility of utilising machine learning models to predict future cluster machine workloads



# Cluster Trace Analysis



# Initial Observations

Based the analysis on prior observations and research of cluster usage trends:

- Resource usage stability
- Job durations
- Repeated job schedule patterns
- Machine attributes






# Python Analyzer

- Built a script to assist with extensive analysis of trace
- Collecting data samples from sections of the trace
  - Sampled from starting timestamp with specified duration
  - Randomly sample a number of machines to
  - 5 minute sampling intervals of job resource usage of selected machines
- Produce statistics and graphs for analytical purposes



# Python Analyzer

Example visualisation of 4 different machine workload timeseries data





# Platform ID Analysis


- Machines grouped by 4 different platform IDs
- Describes machine builds and specifications

| Platform ID Specifications                   |                     |                  |                                     |
|--|---------------------|------------------|-------------------------------------|
| Platform ID (hashed)                         | CPU cores available | Memory available | Amount of machines with platform ID |
| HofLGzk10r/8Ildj2+Lqv0UGGvY82NLoni8+J/YyORU= | 0.5                 | 0.2493 - 0.749   | 11632                               |
| 7OZ0vysYGtB6j9MUHMPzA2Iy7GRzWeJTdXOYCLRKGVg= | 0.25                | 0.2498           | 123                                 |
| GtXakjpd0CD41brK7k/27s3Eby3RpJKy7taB9S8UQRA= | 1                   | 1                | 796                                 |
| JQ1tVQBMHBAlISU1gUNXk2powhYumYA+4cB3KzU29l8= | n/a                 | n/a              | 32                                  |



# Platform ID Analysis

Visualisation of overall machine usage grouped by their platform ID





# Correlation Analysis

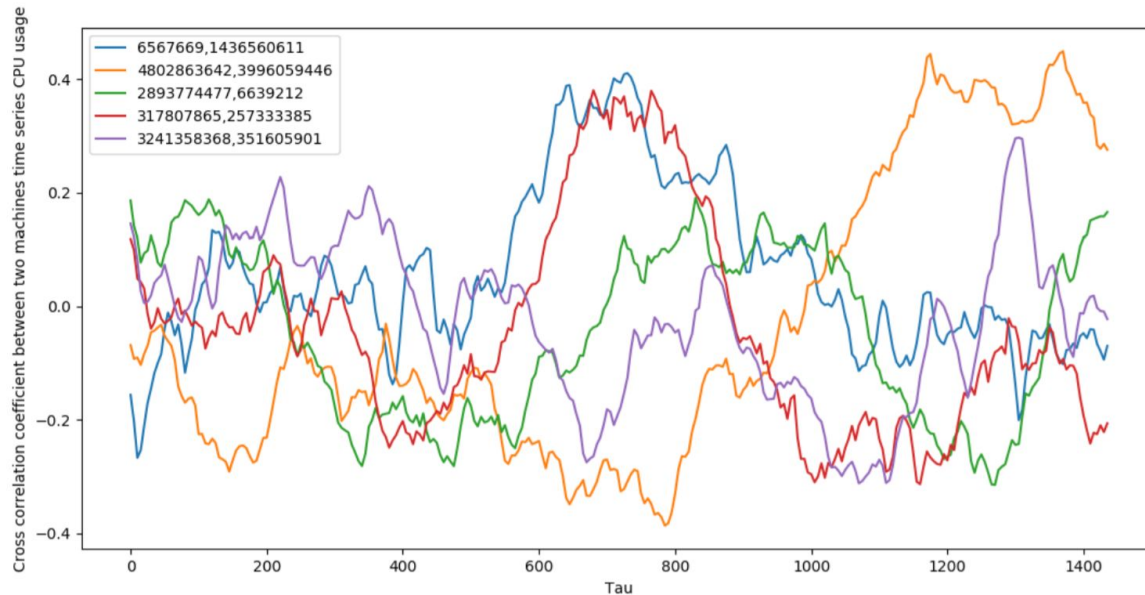
- Analyse the relationship between machine subsets
- Evaluate the strength of the relationship using correlation coefficient

$$r = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_{i=1}^n (x_i - \bar{x})^2 (y_i - \bar{y})^2}}$$



# Correlation Analysis


Example of correlation analysis to analyse the extent of spatial correlation between a pair of machine workloads





# Spatial Correlation Analysis


Cross-correlation between pairs of machine timeseries, averaged across 500 random different pairs





# Spatial Correlation Analysis

Correlation between pairs of machine timeseries, averaged across 50 different pairs from sampled from specific platform ID








# Temporal Correlation Analysis

Autocorrelation within an individual machine timeseries, averaged across 1000 different machines





# Trace Analysis Conclusions

- Overall weak spatial correlation between machine pairs timeseries despite machine attribute characteristics
- Difficult to characterise machine subgroups due to complexity of trace
- Strong temporal correlation within individual machine timeseries



# Forecasting Model





# Neural Network Model Overview

- Implementation of recurrent neural network model which includes:
  - 1D Convolutional layer
  - LSTM unit

| Hyperparameters for RNN                  |                          |
|--|--------------------------|
| Parameter                                | Value                    |
| Number of epochs                         | 500                      |
| Loss function                            | Mean squared error (MSE) |
| Batch size                               | 60                       |
| Input sequence length (timesteps)        | 5                        |
| Sampling rate                            | 1 minute                 |
| Number of hidden LSTM layers             | 2                        |
| Number of neurons in hidden layers       | h1: 128, h2: 64          |
| 1D convolutional kernel size             | 4                        |
| 1D MaxPool kernel size                   | 2                        |
| Output activation function               | Linear                   |
| Optimizer                                | Adam                     |
| Learning rate                            | 0.00059                  |
| L2 regularization penalty (weight decay) | 0.4                      |



# Datasets

Compiled data from Python analyzer, sampled 24 hours of machine workload series data within the cluster trace from

- **Training set** - sampled from 12 hours of data
- **Validation set** - sampled from 2 hours of data




# Hyperparameter Tuning Approach

- Tuning and testing on one set of workload samples from **one specific machine** workload
- Increased sampling rate of dataset (1 min intervals)
- Batch training samples
- Experimentation of various different hyperparameter values



# Data Preprocessing for Time Series Features

- Each sample fed into the LSTM requires a set sequence of timesteps
- The model uses the sequence as learning features to predict values
- Use sequence length as a model hyperparameter that can be varied





# Data Preprocessing for Time Series Features

Example of modification of dataset [ 23, 55, 81, 93, 89, 82, 63, 47, ... ] with timestep sequence length of 3


$$\begin{bmatrix} 23 & 55 & 81 \\ 55 & 81 & 93 \\ 81 & 93 & 89 \\ 93 & 89 & 82 \\ 89 & 82 & 63 \end{bmatrix}$$

Output targets of [ 93, 89, 82, 63, 47 ]





# Timestep Feature Length






# LSTM architecture

| Comparing MSE values for different LSTM architectures |               |               |                |               |
|---|---------------|---------------|----------------|---------------|
| LSTM architecture (col.)<br>Machine ID (row)          | (64, 32)      | (128, 64)     | (128, 128, 64) | (128, 64, 64) |
| 350588109   | 143.163       | 141.965       | 141.673        | 141.701       |
| 348805021   | 289.493       | 278.778       | 248.639        | 262.673       |
| 4820285492  | 34.659        | 19.203        | 19.199         | 20.984        |
| 1436333635  | 149.488       | 97.981        | 153.993        | 126.080       |
| 3338000908  | 123.556       | 58.838        | 168.168        | 168.725       |
| <b>Average</b>  | <b>148.07</b> | <b>119.35</b> | <b>146.33</b>  | <b>144.03</b> |



# LSTM architecture




Visualisation of validation results over different LSTM architectures




# 1D Convolutional Layer

Can assist with processing a representation for complex scheduling features of machine workload timeseries



(a) Without 1D conv. layer,  $MSE = 142.088$



(b) With 1D conv. layer,  $MSE = 119.557$



# Evaluation





# Testing Dataset

- Trained on initial training dataset
- Collected from a different starting timestamp but on the same machine workload series
- Predicting future resource usage on unseen data
- Time sample period of 24 hours



# Evaluation Metrics

- Root Mean Squared Error (RMSE)

$$\sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$$

- Mean Absolute Error (MAE)

$$\frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$$




# Results

| Evaluation results of sampled machine time series |            |  |        |        |
|---|------------|--|--------|--------|
| Index ref.  | Machine ID | Platform ID                                      | RMSE   | MAE    |
| 1   | 348805021  | HofLGzk10r/8Ildj2+Lqv0UG<br>GvY82Nloni8+J/Yy0RU= | 29.501 | 24.763 |
| 2   | 350588109  | HofLGzk10r/8Ildj2+Lqv0UG<br>GvY82Nloni8+J/Yy0RU= | 14.480 | 11.506 |
| 3   | 1436333635 | HofLGzk10r/8Ildj2+Lqv0UG<br>GvY82Nloni8+J/Yy0RU= | 29.501 | 22.770 |
| 4   | 3338000908 | 70Z0vysYGtB6j9MUHMPzA2Iy<br>7GRzWeJTdX0YCLRKGVg= | 20.697 | 16.527 |
| 5   | 1390835522 | 70Z0vysYGtB6j9MUHMPzA2Iy<br>7GRzWeJTdX0YCLRKGVg= | 26.306 | 20.914 |
| 6   | 1391018274 | 70Z0vysYGtB6j9MUHMPzA2Iy<br>7GRzWeJTdX0YCLRKGVg= | 28.891 | 23.176 |
| 7   | 4820285492 | GtXakjpd0CD41brK7k/27s3E<br>by3RpJKy7taB9S8UQRA= | 13.195 | 10.63  |
| 8   | 5015788232 | GtXakjpd0CD41brK7k/27s3E<br>by3RpJKy7taB9S8UQRA= | 12.428 | 10.187 |
| 9   | 4874102959 | GtXakjpd0CD41brK7k/27s3E<br>by3RpJKy7taB9S8UQRA= | 13.797 | 10.937 |





# Results



Visualisation of results for machine workload predictions




# Results

| Additional evaluation metrics based on platform ID |              |              |             |
|--|--------------|--------------|-------------|
| Platform ID  | Avg.<br>RMSE | Avg.<br>MAE  | Difference  |
| HofLGzk10r/8Ildj2+Lqv0UG<br>GvY82NLoni8+J/Yy0RU=   | 24.49        | 19.68        | 4.81        |
| 70Z0vysYGtB6j9MUHMPzA2Iy<br>7GRzWeJTdX0YCLRKGVg=   | 25.30        | 20.21        | 5.09        |
| GtXakjpd0CD41brK7k/27s3E<br>by3RpJKy7taB9S8UQRA=   | 13.14        | 10.58        | 2.56        |
| <b>Overall</b>                                     | <b>20.98</b> | <b>16.82</b> | <b>4.16</b> |




# Evaluating forecasting results



Machine ID: 348805021, RMSE: 29.501, MAE: 24.763




# Evaluating forecasting results



Machine ID: 3338000908, RMSE: 20.697, MAE: 16.527



# Evaluating forecasting results



Machine ID: 4820285492, RMSE: 13.195, MAE: 10.63



## Conclusions from results

- Model generally performs well over the tested machine workloads
- Performs better on machines from platform ID with highest specifications
- Difficult time adapting to spikes and fluctuations within workload series



## Further Improvements

- Prediction with live results to improve our forecasting model
- Use of grouped machine workload timeseries data to train new forecasting model
- Prediction of other cluster usage statistics such as machine failure likelihood at a certain time

# Thank you for listening.

Any questions?

